

The Top-Degree Mass of Unit-Distance Graphs and the Price of Node Privacy

Anonymous author

Anonymous affiliation

Abstract

Erdős asked how many pairs among n points in the plane can be at distance exactly 1. We study an unbalanced version of this question and show that it determines the cost of a basic task in data privacy. Let $\Lambda_k(n)$ be the largest possible sum of the k largest degrees in the unit-distance graph of at most n points. We prove that the smallest expected error with which the number of unit-distance pairs can be released under ε -node differential privacy is $\Theta(\Lambda_{\lceil 1/\varepsilon \rceil}(n))$, with absolute constants, for every $\varepsilon \in (0, 1]$. We also show that for $k \leq n/2$, $\Lambda_k(n)$ equals, up to a factor of 2, the maximum number of incidences between n points and k unit circles with arbitrary centres. We then compare norms on the plane. $\Lambda_k(n) = \Theta(kn)$ if the unit circle contains a segment; $\Lambda_k(n) = O(n^{2/3}k^{2/3} + n)$ if the norm is strictly convex, and a norm found by Valtr attains this bound for every k ; by a theorem of Alon, Bucić and Sauerermann, $\Lambda_k(n) = \Theta(n)$ up to a logarithmic factor for all norms outside a meagre set. For the Euclidean norm the optimal error is $\Theta(n)$ when $\varepsilon \geq n^{-1/2}$, a factor $1/\varepsilon$ below the standard Laplace mechanism. It is polynomially larger than for these generic norms if and only if Erdős's unit-distance conjecture fails, which was shown in 2026. Between these ranges the Euclidean question is an open problem on unbalanced unit-circle incidences.

2012 ACM Subject Classification Theory of computation $\rightarrow$ Computational geometry; Security and privacy $\rightarrow$ Data anonymization and sanitization; Mathematics of computing $\rightarrow$ Combinatorics

Keywords and phrases unit distances, incidences, normed planes, differential privacy, node privacy

Digital Object Identifier 10.4230/LIPIcs...

1 Introduction

Let P be a finite set of points in the plane and let $G(P)$ be its *unit-distance graph*: two points are adjacent if their distance is exactly 1. Erdős [13] asked for the maximum number $u(n)$ of edges of $G(P)$ over n -point sets. The upper bound $u(n) = O(n^{4/3})$ of Spencer, Szemerédi and Trotter [26] has not been improved. Erdős conjectured $u(n) = n^{1+o(1)}$; in 2026 this conjecture was disproved [20, 1], and Sawin [24] gave sets with more than $n^{1.014}$ unit distances for arbitrarily large n .

This paper studies a refinement of $u(n)$ that measures how unevenly unit distances can be distributed among the points. For a finite set P and an integer $k \geq 1$, let $S_k(P)$ be the sum of the k largest degrees of $G(P)$ (padded with zeros if $k > |P|$), and define the *top-degree mass*

$$\Lambda_k(n) = \max\{S_k(P) : |P| \leq n\}.$$

In this definition the k points whose degrees are summed belong to P ; the maximum is over all sets of at most n distinct points. Thus $\Lambda_1(n) = n - 1$ for $n \geq 2$ (a point together with $n - 1$ points on its unit circle), and $\Lambda_k(n) = 2u(n)$ for $k \geq n$. In between, $\Lambda_k(n)$ asks how many unit distances k points can see.

The privacy problem. Our motivation comes from differential privacy. Suppose each point of P is the location of one person, and we wish to publish the number $U(P)$ of unit-distance pairs. *Node differential privacy* requires that adding or removing one person, together with all pairs that contain this person, changes the distribution of the published value by at most a



© Anonymous author(s);

licensed under Creative Commons License CC-BY 4.0

Leibniz International Proceedings in Informatics

LIPICS Schloss Dagstuhl – Leibniz-Zentrum für Informatik, Dagstuhl Publishing, Germany

XX:2 Top-Degree Mass and Node Privacy

factor e^ε (see Section 3 for a self-contained introduction). One point can form a unit-distance pair with all others, so the standard Laplace mechanism must add noise of order n/ε . We ask for the *minimax error*

$$\text{err}^*(\varepsilon, n) = \inf_M \sup_{|P| \leq n} \mathbb{E} |M(P) - U(P)|,$$

where M ranges over ε -node-private algorithms and n is a public bound on the number of points. Our first result says that, up to constant factors, this privacy question is the geometric question above.

► **Theorem 1 (Reduction).** *For every norm on $\mathbb{R}^2$, every $n \geq 1$ and every $\varepsilon \in (0, 1]$, with $k = \lceil 1/\varepsilon \rceil$,*

$$\frac{1}{8(1+e)} \Lambda_k(n) \leq \text{err}^*(\varepsilon, n) \leq 2 \Lambda_k(n).$$

The upper bound is attained by an algorithm that is ε -node-private on all finite point sets. The same holds, with Λ_k defined accordingly, for every family of point sets closed under deleting points and every graph that is determined by pairs of points (for instance, any distance-bin graph).

The ingredients are known in the privacy literature: a bounded-degree surrogate for the edge count due to Kasiviswanathan, Nissim, Raskhodnikova and Smith [19], and a lower bound by deleting and re-inserting high-degree vertices, which is closely related to the down-sensitivity analysis of Dong, Fang, Yi, Tao and Machanavajjhala [9]. We do not claim these tools as new. What Theorem 1 adds is that, over a family closed under deletion, the worst-case cost is given by a single extremal function. This turns the privacy question into a question about point sets, which is the subject of the rest of the paper.

Geometry of the top-degree mass. The k largest degrees of $G(P)$ count incidences between P and the unit circles centred at k points of P . We compare this with the same count for circles whose centres may lie anywhere. Let $I(n, k)$ be the maximum, over sets B of at most n points and sets R of exactly k points of $\mathbb{R}^2$, of the number of pairs $(b, r) \in B \times R$ with $\|b - r\| = 1$. Equivalently, $I(n, k)$ is the maximum number of incidences between at most n points and k distinct unit circles $r + \mathbb{S}$, $r \in R$; distinct centres give distinct circles, and the sets B and R may intersect.

► **Theorem 2 (Top-degree mass and incidences).** *For every norm on $\mathbb{R}^2$ and all integers $n \geq 2$ and $1 \leq k \leq n$,*

$$\max \left\{ n - 1, \frac{2k u(n)}{n}, I(n - k, k) \right\} \leq \Lambda_k(n) \leq I(n, k),$$

where $u(n)$ is the maximum number of unit distances among n points in this norm and $I(0, k) = 0$. If $k \leq n/2$, then also $I(n, k) \leq 2\Lambda_k(n)$.

Which values can Λ_k take? We compare norms on $\mathbb{R}^2$. A norm is *strictly convex* if its unit circle contains no segment. Valtr [30] found a strictly convex norm in which n points can span $\Omega(n^{4/3})$ unit distances; we use it in the form $\|(x, y)\|_V = |y| + \sqrt{x^2 + y^2}$, whose unit ball $\{|y| \leq (1 - x^2)/2\}$ is a linear image of the ball $\{|y| + x^2 \leq 1\}$ (a linear change of coordinates does not change Λ_k). We show that this norm is extremal for every k , not only for $k = n$. Alon, Bucić and Sauermaun [2] consider the space of all norms on $\mathbb{R}^d$ and prove

that, outside a meagre subset of this space (“for almost all norms” in the sense of Baire category), every n points span at most $\frac{d}{2}n \log_2 n$ unit distances. We call the norms outside such a meagre set *generic*. This is a topological notion; it does not concern any particular norm, such as the Euclidean one.

► **Theorem 3** (Top-degree mass across norms). *There are absolute constants $C, c > 0$ such that for all $n \geq 8$ and $1 \leq k \leq n$:*

1. *If the unit circle contains a segment, then $\lfloor n/2 \rfloor \min\{k, n-1\} \leq \Lambda_k(n) \leq (n-1)k$.*
2. *If the norm is strictly convex, then $\Lambda_k(n) \leq C(n^{2/3}k^{2/3} + n)$.*
3. *For Valtr’s norm $\|\cdot\|_V$, we have $\Lambda_k(n) \geq c(n^{2/3}k^{2/3} + n)$.*
4. *For every norm outside a meagre set of norms (in the sense of [2]), $n-1 \leq \Lambda_k(n) \leq 2n \log_2 n$.*

Every norm is either strictly convex or has a segment on its unit circle, so parts (1) and (2) give an upper bound on $\Lambda_k(n)$ for every norm. Part (1) is tight for every norm with a segment. Part (3) shows only that the bound in part (2) is the best one that follows from strict convexity alone: an individual strictly convex norm, such as the Euclidean norm, may have a much smaller $\Lambda_k(n)$. Part (4) concerns generic norms only.

► **Corollary 4** (Privacy cost across norms). *Let $n \geq 8$ and $1/(n-1) \leq \varepsilon \leq 1$. If the unit circle contains a segment, then $\text{err}^*(\varepsilon, n) = \Theta(n/\varepsilon)$, so the Laplace mechanism is optimal. For every strictly convex norm, $\text{err}^*(\varepsilon, n) = O(n + n^{2/3}\varepsilon^{-2/3})$, and for Valtr’s norm also $\text{err}^*(\varepsilon, n) = \Omega(n + n^{2/3}\varepsilon^{-2/3})$. For every generic norm, $(n-1)/(8(1+e)) \leq \text{err}^*(\varepsilon, n) \leq 4n \log_2 n$; this holds for all $n \geq 2$ and $\varepsilon \in (0, 1]$.*

The Euclidean plane. The Euclidean norm is strictly convex, so its privacy cost lies below Valtr’s curve. The averaging bound in Theorem 2 brings in $u(n)$.

► **Theorem 5** (Euclidean plane). *There are absolute constants $C, c > 0$ such that for the Euclidean norm, all $n \geq 8$ and all $\varepsilon \in [1/n, 1]$,*

$$c \left(n + \frac{u(n)}{\varepsilon n} \right) \leq \text{err}^*(\varepsilon, n) \leq C \left(n + n^{2/3}\varepsilon^{-2/3} \right).$$


In particular $\text{err}^(\varepsilon, n) = \Theta(n)$ for $\varepsilon \geq n^{-1/2}$. Moreover, the following are equivalent:*

1. *$u(n) \geq n^{1+\delta}$ for some $\delta > 0$ and infinitely many n ;*
2. *there are $\delta' > 0$ and privacy levels $\varepsilon_n \in (0, 1]$ such that $\text{err}^*(\varepsilon_n, n) \geq n^{1+\delta'}$ for infinitely many n .*

By the 2026 results, both hold: for every $\eta \in [0, 0.014)$ there are infinitely many n with $\text{err}^(n^{-1+\eta}, n) \geq cn^{1.014-\eta}$, while for every generic norm the error is at most $4n \log_2 n$ at every ε .*

Figure 1 shows the resulting picture. For the Euclidean norm, the region between the two bounds of Theorem 5 is open. By Theorem 2, closing it is the same as determining $I(n, k)$ for unit circles when $\sqrt{n} \leq k \leq n$, an unbalanced version of the unit-distance problem. At $k = n$ this is the unit-distance problem itself; Theorem 3(3) shows that any improvement over the incidence bound must use properties of the Euclidean norm that Valtr’s norm does not have. We state the question in Section 8.

Summary and provenance. Table 1 lists the results and says which parts are new. We regard the geometric results (Theorems 2, 3, and 5) as the main contribution. Theorem 1 is the bridge; Section 7 adds an algorithm that does not need to know Λ_k , at a logarithmic cost.



■ **Figure 1** Exponent of the optimal error against $\log_n(1/\varepsilon)$. Shaded: the range allowed for the Euclidean norm by Theorem 5; its lower edge is drawn with $\delta = 0.06$, while the best known value is $\delta \geq 0.014$.

Result	Statement	Status
Theorem 1	$\text{err}^*(\varepsilon, n) = \Theta(\Lambda_{\lceil 1/\varepsilon \rceil}(n))$	known tools; new form
Theorem 2	$\Lambda_k(n)$ and $I(n, k)$ agree up to a factor 2 ($k \leq n/2$)	new (short)
Theorem 3(1),(4)	segment norms; generic norms	new; uses [2]
Theorem 3(2)	strictly convex norms	known incidence bound
Theorem 3(3)	Valtr's norm is extremal for every k	new; extends [30]
Theorem 5	Euclidean window; separation iff $u(n) \geq n^{1+\Omega(1)}$	new connection
Theorem 13	algorithm without knowledge of Λ_k	known tools [19, 23]

■ **Table 1** Results and their provenance.

125 **Related work.** *Unit distances and incidences.* Besides the bounds above, the unit-distance
126 problem is surveyed by Brass, Moser and Pach [5]. The $O(n^{4/3})$ bound follows from incidence
127 bounds for points and curves [29, 7, 21], and Székely's crossing-number argument [28] gives it
128 for any family of curves in which two curves meet in at most two points and two points lie on
129 at most two curves. For normed planes, Swanepoel's survey [27] describes the known results,
130 including Valtr's construction. After the disproof of Erdős's conjecture [20, 1], Sawin [24]
131 gave the exponent 1.014 and Emmerich [12] optimised explicit finite certificates. We use
132 these results only through Theorem 2; we do not use their structure.

133 *Node privacy.* Node differential privacy for graphs was studied by Blocki, Blum, Datta and
134 Sheffet [4] and by Kasiviswanathan, Nissim, Raskhodnikova and Smith [19], who introduced
135 the bounded-degree flow statistic we use. For a family closed under deletion, the restricted
136 sensitivity of Blocki, Blum, Datta and Sheffet is $\Lambda_1(n)$ in our notation, which leads to error
137 about $\Lambda_1(n)/\varepsilon$. Theorem 1 replaces this by $\Lambda_{\lceil 1/\varepsilon \rceil}(n)$, which is never larger; for strictly
138 convex norms it is smaller by a factor of order at least $\varepsilon^{-2/3}$. Raskhodnikova and Smith [23]
139 developed Lipschitz extensions and the generalized exponential mechanism used in Section 7.
140 Instance-dependent guarantees in terms of down sensitivity appear in work of Asi and
141 Duchi [3], Fang, Dong and Yi [14], and Dong, Fang, Yi, Tao and Machanavajjhala [9]; see
142 also Cummings and Durfee [8]. Recent node-private results include counting connected
143 components [18], continual observation [17] and graphon estimation [6]. Raskhodnikova,
144 Smith, Wagaman and Zavyalov [22] study node privacy in the *local* model, where each person
145 randomizes their own edge list; we work in the central model, where a trusted curator holds

the data. We are not aware of earlier work that links node privacy to incidence geometry.

2 Overview of the proofs

Why the privacy cost is $\Lambda_{\lceil 1/\varepsilon \rceil}$. The upper bound uses one public threshold. The surrogate g_D of [19] changes by at most D when a point is added or removed, and it undercounts $U(P)$ by at most the degree tail $\text{tail}_D(P) = \sum_p (\deg_P(p) - D)_+$. Set $k = \lceil 1/\varepsilon \rceil$ and $D = \Lambda_k(n)/k$. The vertices of degree above D form a set S , and $\text{tail}_D(P) \leq S_{|S|}(P) - |S|D \leq \Lambda_{|S|}(n) - |S|D$. The average $\Lambda_j(n)/j$ cannot increase with j , so this is at most $\Lambda_k(n)$ for every value of $|S|$. Laplace noise of scale $D/\varepsilon \leq \Lambda_k(n)$ completes the bound.

For the lower bound, take a set P^* with $S_k(P^*) = \Lambda_k(n)$, where now $k \approx 1/\varepsilon$, and remove its k points of largest degree. Re-inserting them one at a time is a path of $k \leq 1/\varepsilon$ neighbouring inputs along which the count grows by at least $\Lambda_k(n)/2$. A private algorithm cannot distinguish the two ends of such a path well, so it errs by $\Omega(\Lambda_k(n))$ on one of them.

Why Λ_k is an incidence count. If C is the set of k points of largest degree in P , then $S_k(P)$ is the number of incidences between P and the unit circles centred at points of C . Conversely, given n points and k unit circles, adding the k centres to the point set turns incidences into degrees. Splitting the points into two halves keeps the total size below n at a loss of a factor 2.

Why the norms differ. If the unit circle contains a segment, two parallel rows of points give a complete bipartite unit-distance graph, so $\Lambda_k(n) = \Theta(kn)$. If the norm is strictly convex, two distinct unit circles meet in at most two points and two points lie on at most two unit circles; the crossing-number argument then gives $I(n, k) = O(n^{2/3}k^{2/3} + n + k)$. For Valtr's norm, the upper half of each unit circle is a parabolic arc, and the map $(x, y) \mapsto (x, y + x^2/2)$ sends all these arcs to segments of lines. The standard grid construction for point-line incidences, at the right aspect ratio, then gives an exact count of unit distances (Lemma 12) and $\Lambda_k(n) = \Omega(n^{2/3}k^{2/3})$ for every $\sqrt{n} \leq k \leq n$. For generic norms the bound of [2] on $u(n)$ caps $\Lambda_k(n) \leq 2u(n)$.

Why the 2026 result enters. The averaging bound $\Lambda_k(n) \geq 2k u(n)/n$ holds because the k largest of n degrees have at least the average degree. With $k = \lceil 1/\varepsilon \rceil$ close to n , Theorem 1 turns a superlinear $u(n)$ into a superlinear privacy cost. Conversely $\Lambda_k(n) \leq \Lambda_n(n) = 2u(n)$, so no other source of superlinear cost exists.

3 Preliminaries

Geometry. Fix a norm $\|\cdot\|$ on $\mathbb{R}^2$ with unit circle $\mathbb{S} = \{x : \|x\| = 1\}$; a *unit circle* is a translate $c + \mathbb{S}$. For a finite set P , $\deg_P(p)$ is the number of $q \in P$ with $\|p - q\| = 1$, and $U(P)$ is the number of edges of $G(P)$. For a finite family Γ of curves, $I(P, \Gamma)$ is the number of pairs $(p, \gamma) \in P \times \Gamma$ with $p \in \gamma$. Since $\mathbb{S} = -\mathbb{S}$, a point p lies on $c + \mathbb{S}$ exactly when $\|p - c\| = 1$. We write $\mathcal{P}_n$ for the family of point sets with at most n points. The functions $\Lambda_k(n)$ and $I(n, k)$ are non-decreasing in both arguments, and $\Lambda_k(n) = \Lambda_n(n) = 2u(n)$ for $k \geq n$.

Node differential privacy, for geometers. An *algorithm* M maps each finite point set P to a random real number $M(P)$. Two point sets are *node neighbours* if one is obtained from the

XX:6 Top-Degree Mass and Node Privacy

other by adding one point. The word “node” refers to the graph $G(P)$: adding a point adds a vertex together with all its edges.

► **Definition 6.** Let $\varepsilon > 0$ and let $\mathcal{F}$ be a family of finite point sets. An algorithm M is ε -node-private on $\mathcal{F}$ if for all node neighbours $P, Q \in \mathcal{F}$ and every measurable $T \subseteq \mathbb{R}$, $\Pr[M(P) \in T] \leq e^\varepsilon \Pr[M(Q) \in T]$.

The definition says that the output reveals little about whether any single person is present. Smaller ε means stronger privacy; the range $\varepsilon \leq 1$ is the usual one. We use three standard facts [10, 11]. (i) *Laplace mechanism*: if $|f(P) - f(Q)| \leq \Delta$ for all node neighbours, then $f(P) + Z$ with Z drawn from the Laplace density $\frac{\varepsilon}{2\Delta} e^{-\varepsilon|z|/\Delta}$ is ε -node-private, and $\mathbb{E}|Z| = \Delta/\varepsilon$. For $f = U$ one has $\Delta = n - 1$, which gives error $(n - 1)/\varepsilon$; we call this the *standard Laplace mechanism*. (ii) *Groups*: if P and Q are joined by a path of h node neighbours in $\mathcal{F}$, then $\Pr[M(P) \in T] \leq e^{h\varepsilon} \Pr[M(Q) \in T]$. (iii) The following two-point bound, which we prove for completeness.

► **Lemma 7.** Let M be ε -node-private on $\mathcal{F}$, and let $P, Q \in \mathcal{F}$ be joined by a path of h node neighbours in $\mathcal{F}$ with $h\varepsilon \leq 1$. If $U(Q) - U(P) = \Delta > 0$, then $\mathbb{E}|M(P) - U(P)|$ or $\mathbb{E}|M(Q) - U(Q)|$ is at least $\Delta/(2(1 + e))$.

Proof. Let α be the larger of the two expected errors and $T = [U(P) + \Delta/2, \infty)$. By Markov’s inequality $\Pr[M(P) \in T] \leq 2\alpha/\Delta$ and $\Pr[M(Q) \notin T] \leq 2\alpha/\Delta$. By (ii), $1 - 2\alpha/\Delta \leq \Pr[M(Q) \in T] \leq e \cdot 2\alpha/\Delta$, so $\alpha \geq \Delta/(2(1 + e))$. ◀

4 The reduction

► **Lemma 8.** For every n , the sequence $k \mapsto \Lambda_k(n)/k$ is non-increasing. Hence $\Lambda_j(n) \leq (j/k)\Lambda_k(n)$ for $j \geq k$.

Proof. For a fixed P , $S_k(P)/k$ is the average of the k largest degrees (with zero padding), which does not increase when k grows. A pointwise maximum of non-increasing sequences is non-increasing. ◀

For a graph G and a real number $D \geq 0$, let

$$g_D(G) = \max \left\{ \sum_{e \in E(G)} x_e : 0 \leq x_e \leq 1, \sum_{e \ni v} x_e \leq D \text{ for every vertex } v \right\}.$$

It equals one half of the maximum-flow value in the network of Kasiviswanathan, Nissim, Raskhodnikova and Smith [19] (Proposition 14 in Appendix A). We write $g_D(P)$ for $g_D(G(P))$.

► **Lemma 9.** For every real $D \geq 0$: (a) if P, Q are node neighbours then $|g_D(P) - g_D(Q)| \leq D$; (b) $U(P) - \text{tail}_D(P) \leq g_D(P) \leq U(P)$ for every P .

Proof. (a) Let $Q = P \cup \{q\}$. An optimal solution for Q , restricted to the edges not containing q , is feasible for P and loses at most $\sum_{e \ni q} x_e \leq D$. An optimal solution for P , extended by zeros, is feasible for Q . (b) The upper bound holds since $x_e \leq 1$. For the lower bound put $x_e = \min\{1, D/\deg u, D/\deg v\}$ for $e = uv$. At each vertex v the sum is at most $\deg v \cdot \min\{1, D/\deg v\} \leq D$. Each edge loses $1 - x_e \leq (1 - D/\deg u)_+ + (1 - D/\deg v)_+$, and summing over edges gives $\sum_v \deg v (1 - D/\deg v)_+ = \text{tail}_D(P)$. ◀

► **Lemma 10.** For every $P \in \mathcal{P}_n$ and real $D \geq 0$, $\text{tail}_D(P) \leq \max_{0 \leq j \leq n} (\Lambda_j(n) - jD)$, where $\Lambda_0 = 0$.

Proof. Let S be the set of points of degree larger than D and $j = |S|$. Then $\text{tail}_D(P) = \sum_{p \in S} \deg_P(p) - jD \leq S_j(P) - jD \leq \Lambda_j(n) - jD$. $\blacktriangleleft$

The proof below uses only properties (a) and (b) of Lemma 9. Any statistic with these two properties can replace g_D ; for integer D this includes the integral statistic of Appendix A.

► **Theorem 1** (Reduction). *For every norm on $\mathbb{R}^2$, every $n \geq 1$ and every $\varepsilon \in (0, 1]$, with $k = \lceil 1/\varepsilon \rceil$,*

$$\frac{1}{8(1+e)} \Lambda_k(n) \leq \text{err}^*(\varepsilon, n) \leq 2 \Lambda_k(n).$$

The upper bound is attained by an algorithm that is ε -node-private on all finite point sets. The same holds, with Λ_k defined accordingly, for every family of point sets closed under deleting points and every graph that is determined by pairs of points (for instance, any distance-bin graph).

Proof. *Upper bound.* Let $k = \lceil 1/\varepsilon \rceil$, $\lambda = \Lambda_k(n)$ and $D = \lambda/k$. The algorithm outputs $g_D(P) + Z$ with Z Laplace of scale D/ε ; it is ε -node-private on all finite sets by Lemma 9(a). For $j \geq k$, Lemma 8 gives $\Lambda_j(n) - jD \leq 0$, and for $j < k$ we have $\Lambda_j(n) - jD \leq \lambda$. By Lemma 10 and Lemma 9(b), for every $P \in \mathcal{P}_n$,

$$\mathbb{E}|g_D(P) + Z - U(P)| \leq \text{tail}_D(P) + \mathbb{E}|Z| \leq \lambda + \frac{\lambda}{k\varepsilon} \leq 2\lambda.$$

Lower bound. Let $m = \min\{\lceil 1/\varepsilon \rceil, n\} \geq 1$ and choose $P^* \in \mathcal{P}_n$ with $S_m(P^*) = \Lambda_m(n)$. Let C consist of $\min\{m, |P^*|\}$ points of P^* of largest degree and $P_0 = P^* \setminus C$. Adding the points of C to P_0 one at a time gives a path of $|C| \leq m$ node neighbours inside $\mathcal{P}_n$. Every edge counted in $\sum_{c \in C} \deg(c) = \Lambda_m(n)$ is counted at most twice, so $U(P^*) - U(P_0) \geq \Lambda_m(n)/2$. Since $m\varepsilon \leq 1$, Lemma 7 shows that every ε -node-private algorithm has expected error at least $\Lambda_m(n)/(4(1+e))$ on P_0 or on P^* . Finally, if $m = n$ then $\Lambda_k(n) = \Lambda_m(n)$; otherwise $k \leq \lceil 1/\varepsilon \rceil + 1 \leq 2m$, and Lemma 8 gives $\Lambda_k(n) \leq 2\Lambda_m(n)$.

Other families. The upper bound used only Lemma 9, which holds for all graphs. The lower bound used only that P_0 and the intermediate sets belong to the family. Both hold for any family closed under deleting points and any graph in which adjacency of two points depends only on these two points. $\blacktriangleleft$

► **Remark 11.** The algorithm needs the value $\Lambda_k(n)$. If only an upper bound $\lambda \geq \Lambda_k(n)$ is known, the same proof with $D = \lambda/k$ gives error at most 2λ . For the Euclidean norm this yields error $O(n + n^{2/3}\varepsilon^{-2/3})$ from Theorem 3(2). Theorem 13 below removes the need for any bound at a logarithmic cost.

5 Top-degree mass and incidences

► **Theorem 2** (Top-degree mass and incidences). *For every norm on $\mathbb{R}^2$ and all integers $n \geq 2$ and $1 \leq k \leq n$,*

$$\max\left\{n-1, \frac{2k u(n)}{n}, I(n-k, k)\right\} \leq \Lambda_k(n) \leq I(n, k),$$

where $u(n)$ is the maximum number of unit distances among n points in this norm and $I(0, k) = 0$. If $k \leq n/2$, then also $I(n, k) \leq 2\Lambda_k(n)$.

Proof. Lower bounds. Placing $n - 1$ points on the unit circle around a further point gives $S_1 \geq n - 1$. If $|P| = n$ and $U(P) = u(n)$, the average degree is $2u(n)/n$, and the k largest degrees have at least this average, so $S_k(P) \geq 2ku(n)/n$.

$\Lambda_k(n) \leq I(n, k)$. Let $P \in \mathcal{P}_n$ and let C be a set of $\min\{k, |P|\}$ points of P of largest degree. Then $S_k(P) = \sum_{c \in C} |P \cap (c + \mathbb{S})|$. If $|C| < k$, add $k - |C|$ further centres far from P ; this does not change the count, so $S_k(P) \leq I(n, k)$.

$I(n - k, k) \leq \Lambda_k(n)$. Let B have at most $n - k$ points and let R be a set of k centres. The set $Q = B \cup R$ has at most n points, and each $r \in R$ satisfies $\deg_Q(r) \geq |B \cap (r + \mathbb{S})|$, because $r \notin r + \mathbb{S}$. Since $S_k(Q)$ is at least the sum of the degrees of any k points of Q , $\Lambda_k(n) \geq S_k(Q) \geq \sum_{r \in R} |B \cap (r + \mathbb{S})|$.

$I(n, k) \leq 2\Lambda_k(n)$ for $k \leq n/2$. Let B be a set of at most n points, R a set of k centres, and $\Gamma_R = \{r + \mathbb{S}\}_{r \in R}$. Split B into two parts of at most $\lceil n/2 \rceil$ points each, and let B' be a part with $I(B', \Gamma_R) \geq I(B, \Gamma_R)/2$. Since $|B'| \leq \lceil n/2 \rceil \leq n - k$, the previous paragraph gives $\Lambda_k(n) \geq I(B', \Gamma_R) \geq I(B, \Gamma_R)/2$. $\blacktriangleleft$

Lemma 10 also bounds the degree tail of every set by the top-degree mass. With Theorem 3(2) this gives, for every strictly convex norm,

$$\text{tail}_D(P) \leq \max_j (C(n^{2/3}j^{2/3} + n) - jD) = O(n^2/D^2 + n) \quad (P \in \mathcal{P}_n, D \geq 1). \quad (1)$$

Both terms are needed. A star with $2D$ leaves, repeated $\lfloor n/(2D + 1) \rfloor$ times, has tail $\Omega(n)$. Since $\text{tail}_D(P) \geq S_k(P) - kD$ for every k , Theorem 3(3) with $k \asymp n^2/D^3$ shows that for Valtr's norm the first term is attained when $n^{1/3} \lesssim D \lesssim n^{1/2}$. We note that (1) applies the incidence bound once, to the circles centred at the high-degree points; summing a bound on the number of j -rich points over all $j > D$ would lose a logarithmic factor.


6 Norms on the plane

► **Theorem 3** (Top-degree mass across norms). *There are absolute constants $C, c > 0$ such that for all $n \geq 8$ and $1 \leq k \leq n$:*

1. *If the unit circle contains a segment, then $\lfloor n/2 \rfloor \min\{k, n - 1\} \leq \Lambda_k(n) \leq (n - 1)k$.*
2. *If the norm is strictly convex, then $\Lambda_k(n) \leq C(n^{2/3}k^{2/3} + n)$.*
3. *For Valtr's norm $\|\cdot\|_V$, we have $\Lambda_k(n) \geq c(n^{2/3}k^{2/3} + n)$.*
4. *For every norm outside a meagre set of norms (in the sense of [2]), $n - 1 \leq \Lambda_k(n) \leq 2n \log_2 n$.*

Proof of (1). Let the segment $\{v + tw : |t| \leq 1\}$ lie on $\mathbb{S}$. Put $m = \lfloor n/2 \rfloor$, $a_i = (i/m)w$ and $b_j = v + (j/m)w$ for $0 \leq i, j < m$. Then $b_j - a_i = v + \frac{j-i}{m}w$ lies on the segment, so all m^2 pairs $\{a_i, b_j\}$ are at unit distance. (The points are distinct, since a segment on $\mathbb{S}$ does not lie on a line through the origin.) Every point has degree at least m , and $2m \geq n - 1$, so $S_k \geq m \min\{k, n - 1\}$. The upper bound holds because every degree is at most $n - 1$. $\blacktriangleleft$

Proof of (2). In a strictly convex norm, two distinct unit circles meet in at most two points, and two distinct points lie on at most two common unit circles (Lemma 18 in Appendix C). For such families, Székely's crossing-number argument [28], or the theorem of Pach and Sharir [21], gives $I(P, \Gamma) \leq C_0(|P|^{2/3}|\Gamma|^{2/3} + |P| + |\Gamma|)$; we include the argument in Appendix C. By Theorem 2, $\Lambda_k(n) \leq I(n, k) \leq C_0(n^{2/3}k^{2/3} + n + k)$, and $k \leq n$. $\blacktriangleleft$



■ **Figure 2** Left: three unit circles of Valtr's norm, each bounded by two parabolic arcs. Right: the map $\Phi(x, y) = (x, y + x^2/2)$ sends the upper arcs (bold) to line segments whose slope is the first coordinate of the centre. Point sets with many point–line incidences therefore yield many unit distances (Lemma 12).

302 ► **Lemma 12** (Valtr grid). *Let $A, B \geq 1$ be integers. For $1 \leq i \leq A$, $1 \leq j \leq 2AB$, $1 \leq s \leq B$*
 303 *and $1 \leq t \leq AB$ put*

$$304 \quad p_{ij} = \left(\frac{i}{2A}, \frac{j}{4AB} - \frac{i^2}{8A^2} \right), \quad c_{st} = \left(\frac{s}{2B}, \frac{t}{4AB} - \frac{1}{2} + \frac{s^2}{8B^2} \right).$$

305 *The $2A^2B$ points p_{ij} are distinct, the AB^2 centres c_{st} are distinct, no point equals a centre,*
 306 *and $\|p_{ij} - c_{st}\|_V = 1$ holds if and only if $j = si + t$. In particular, exactly A^2B^2 point–centre*
 307 *pairs are at unit distance.*

308 **Proof.** For $d = (d_x, d_y)$ we have $\|d\|_V = 1$ if and only if $\sqrt{d_x^2 + d_y^2} = 1 - |d_y|$, that is, if and
 309 only if $d_x^2 = 1 - 2|d_y|$. Every point has second coordinate at least $\frac{1}{4AB} - \frac{1}{8} > -\frac{1}{8}$, and every
 310 centre has second coordinate at most $\frac{1}{4} - \frac{1}{2} + \frac{1}{8} = -\frac{1}{8}$. Hence no point equals a centre, and for
 311 $d = p_{ij} - c_{st}$ we have $d_y > 0$. Write $x = \frac{i}{2A}$ and $a = \frac{s}{2B}$. Then $1 - 2d_y = x^2 + a^2 - \frac{2j}{4AB} + \frac{2t}{4AB}$
 312 and $d_x^2 = x^2 - 2ax + a^2$, so $\|d\|_V = 1$ if and only if $\frac{j}{4AB} = ax + \frac{t}{4AB}$, that is, $j = si + t$. For
 313 each of the AB^2 pairs (s, t) and each $i \leq A$, the value $si + t$ lies in $[2, 2AB]$, so it is a valid
 314 index j . This gives exactly $A \cdot AB^2$ pairs. Distinctness is clear from the coordinates. ◀

315 In terms of Figure 2: the map Φ sends p_{ij} to the grid point $(\frac{i}{2A}, \frac{j}{4AB})$ and the upper arc
 316 of $c_{st} + \mathbb{S}$ into the line of slope $\frac{s}{2B}$ and intercept $\frac{t}{4AB}$, so the lemma is the standard grid
 317 example for point–line incidences [29]. The lower arcs carry no incidences, since every point
 318 lies above every centre.

319 **Proof of (3).** The norm $\|\cdot\|_V$ is the sum of the Euclidean norm and the seminorm $|y|$, so
 320 it is strictly convex. Let $m = \lfloor n/2 \rfloor$. If $\sqrt{n} \leq k \leq m$, choose $A, B \geq 1$ with $2A^2B \leq m$,
 321 $AB^2 \leq k$ and $A^2B^2 \geq n^{2/3}k^{2/3}/53$ (Appendix C). The set Q of all points and centres of
 322 Lemma 12 has at most $m + k \leq n$ elements, and $S_k(Q)$ is at least the sum of the degrees of
 323 the $AB^2 \leq k$ centres, which is at least A^2B^2 . The cases $k < \sqrt{n}$ (where $\Lambda_k(n) \geq n - 1$) and
 324 $k > m$ (by monotonicity in k) are treated in Appendix C. ◀

325 **Proof of (4).** By [2] with $d = 2$, for every norm outside a meagre set and every n , every n
 326 points span at most $n \log_2 n$ unit distances. Then $\Lambda_k(n) \leq \Lambda_n(n) = 2u(n) \leq 2n \log_2 n$. The
 327 lower bound $n - 1$ holds in every norm. ◀

328 **Proof of Corollary 4.** Let $k = \lceil 1/\varepsilon \rceil \leq n - 1$ and apply Theorem 1. With a segment,
 329 $\Lambda_k(n) \geq \lfloor n/2 \rfloor / \varepsilon$, and the standard Laplace mechanism has error $(n - 1)/\varepsilon$. For strictly
 330 convex norms, (2) gives $2\Lambda_k(n) = O(n + n^{2/3}k^{2/3})$ and $k \leq 2/\varepsilon$; (3) gives the matching lower
 331 bound for Valtr's norm. For generic norms, Theorem 1 and (4) give $(n - 1)/(8(1 + e)) \leq$
 332 $\text{err}^*(\varepsilon, n) \leq 2\Lambda_k(n) \leq 4n \log_2 n$ for every $\varepsilon \in (0, 1]$. ◀

333 ► **Theorem 5** (Euclidean plane). *There are absolute constants $C, c > 0$ such that for the*
 334 *Euclidean norm, all $n \geq 8$ and all $\varepsilon \in [1/n, 1]$,*

$$335 \quad c \left(n + \frac{u(n)}{\varepsilon n} \right) \leq \text{err}^*(\varepsilon, n) \leq C \left(n + n^{2/3} \varepsilon^{-2/3} \right).$$

336 *In particular $\text{err}^*(\varepsilon, n) = \Theta(n)$ for $\varepsilon \geq n^{-1/2}$. Moreover, the following are equivalent:*

- 337 1. $u(n) \geq n^{1+\delta}$ for some $\delta > 0$ and infinitely many n ;
- 338 2. there are $\delta' > 0$ and privacy levels $\varepsilon_n \in (0, 1]$ such that $\text{err}^*(\varepsilon_n, n) \geq n^{1+\delta'}$ for infinitely
- 339 many n .

340 *By the 2026 results, both hold: for every $\eta \in [0, 0.014)$ there are infinitely many n with*
 341 *$\text{err}^*(n^{-1+\eta}, n) \geq c n^{1.014-\eta}$, while for every generic norm the error is at most $4n \log_2 n$ at*
 342 *every ε .*

343 **Proof.** Let $k = \lceil 1/\varepsilon \rceil$, so $1/\varepsilon \leq k \leq \min\{n, 2/\varepsilon\}$. By Theorems 1 and 2, $\text{err}^*(\varepsilon, n) \geq$
 344 $\frac{1}{8(1+\varepsilon)} \max\{n-1, 2u(n)/(\varepsilon n)\}$. The upper bound follows from Theorems 1 and 3(2). For
 345 $\varepsilon \geq n^{-1/2}$ we have $n^{2/3} \varepsilon^{-2/3} \leq n$.

346 For the equivalence, if $u(n) \geq n^{1+\delta}$ then at $\varepsilon = 1/n$ the lower bound is at least $c u(n) \geq$
 347 $c n^{1+\delta}$, which exceeds $n^{1+\delta/2}$ for large n . Conversely, $\text{err}^*(\varepsilon, n) \leq 2\Lambda_k(n) \leq 2\Lambda_n(n) = 4u(n)$
 348 for every ε , so a lower bound $n^{1+\delta'}$ on the error forces $u(n) \geq n^{1+\delta'}/4$. Finally, Sawin [24]
 349 gives $u(n) > n^{1.014}$ for infinitely many n ; for these n and $\varepsilon = n^{-1+\eta}$ the lower bound is at
 350 least $c n^{-\eta} u(n) \geq c n^{1.014-\eta}$. The comparison with generic norms is Corollary 4. ◀

351 **What the separation does and does not say.** The separation in Theorem 5 appears only
 352 for $\varepsilon \leq n^{-1+0.014}$, that is, for very strong privacy. At such ε the optimal error is already
 353 comparable to the count itself, so we do not see it as a recommendation for practice. Its
 354 content is structural: the shape of the Euclidean rate curve near $\varepsilon = 1/n$ is decided exactly
 355 by whether $u(n)$ is polynomially superlinear, and the 2026 results settle this. For larger ε
 356 the curve is decided by $I(n, k)$ with $k < n$, where the gap in Theorem 5 is open.

357 7 Algorithms and other distance graphs

358 **An algorithm that does not need Λ_k .** The algorithm of Theorem 1 uses the unknown
 359 number $\Lambda_k(n)$. The following result, which combines the surrogate of [19] with the generalized
 360 exponential mechanism of Raskhodnikova and Smith [23], adapts to the input instead. It is
 361 not new as a method; we include it because its guarantee has a direct geometric reading.

362 ► **Theorem 13.** *Let $K = 1 + \lceil \log_2 n \rceil$. There is an ε -node-private algorithm such that for*
 363 *every $P \in \mathcal{P}_n$ and $\beta \in (0, \frac{1}{2})$, with probability at least $1 - \beta$, its error is $O(S_t(P) + t)$ with*
 364 *$t = \lceil \log(K/\beta)/\varepsilon \rceil$. Consequently the error is $O(\log(K/\beta) (\Lambda_{\lceil 1/\varepsilon \rceil}(n) + 1/\varepsilon))$.*

365 The algorithm selects a threshold D from $\{1, 2, 4, \dots, 2^{K-1}\}$ privately and then releases
 366 $g_D(P)$ with Laplace noise; the proof is in Appendix B. The bound $S_t(P)$ is the number of
 367 incidences between P and the t richest unit circles centred in P . For instance, if every point
 368 of P has degree at most Δ , the error is $O(\Delta \log(K/\beta)/\varepsilon)$ regardless of n .

369 **Surrogates.** Theorem 1 needs real thresholds, because $\Lambda_k(n)/k$ need not be an integer; this
 370 is one reason to use the fractional statistic g_D . The integral statistic $f_D(G)$, the maximum
 371 number of edges of a subgraph of maximum degree at most D , satisfies the same two
 372 properties for integer D and can be computed by b -matching [25, 15]. With D rounded up
 373 it gives Theorem 1 up to an additive $1/\varepsilon$ (Appendix A). Both statistics are computable in
 374 polynomial time from $G(P)$.

Other distance graphs. Theorem 1 applies to every family of point sets that is closed under deleting points. Two examples are worked out in Appendix E. If the points lie on a grid $\frac{1}{q}\mathbb{Z}^2$ and the target distance r satisfies $qr \in \mathbb{Z}$, every degree is at most the number Δ of representations of $(qr)^2$ as a sum of two squares, and the error is $\Theta(\Delta/\varepsilon)$ whenever $n \geq \lceil 1/\varepsilon \rceil (\Delta + 1)$. If the points are δ -separated and adjacency means a distance in $[r - t, r + t]$, every degree is $O((t + \delta)(r + t + \delta)/\delta^2)$; in this case releasing g_D for D equal to this bound is private on all inputs and exact on separated inputs. Finally, Appendix D treats a class of point sets defined by an upper bound on the number of j -rich points for every j ; this class is not closed under deletion, so Theorem 1 does not apply, and we give separate matching bounds.

8 Open problems

The Euclidean window. For $\alpha \in [\frac{1}{2}, 1]$ let $\gamma(\alpha) = \limsup_{n \rightarrow \infty} \log \Lambda_{\lceil n^\alpha \rceil}(n) / \log n$ for the Euclidean norm, and let $1 + \delta^* = \gamma(1)$ be the unit-distance exponent. Theorems 2 and 3 give

$$\max\{1, \alpha + \delta^*\} \leq \gamma(\alpha) \leq \frac{2+2\alpha}{3},$$

with $\delta^* \geq 0.014$, and for Valtr's norm the upper bound is attained. Is the lower bound the truth? In words: can a few unit circles be much richer than average while the unit-distance graph as a whole stays sparse? By Theorem 2 this is a question about incidences between n points and k unit circles. We are not aware of results on this unbalanced version beyond the bounds above.

Instance-wise bounds. Theorem 1 concerns the worst case over $\mathcal{P}_n$. A local version would compare an algorithm, at each input P , with the best algorithm on a small neighbourhood of P . Deleting points is controlled by $S_t(P)$, as in Theorem 13; adding points brings in the external quantity $\max_{c \notin P} |P \cap (c + \mathbb{S})|$ and its analogues for several inserted points.

Higher dimensions. In $\mathbb{R}^3$, $u(n)$ lies between $n^{4/3} \log \log n$ [5] and $n^{295/197+o(1)}$ [31], so the window is wider and both ends are open. In $\mathbb{R}^d$ with $d \geq 4$, points on two orthogonal circles of radius $1/\sqrt{2}$ form a complete bipartite unit-distance graph, as in Theorem 3(1), and the standard Laplace mechanism is optimal.

Use of generative AI. A generative AI assistant (Claude, Anthropic) was used to draft the text and L^AT_EX source of this version, to propose and check proof arguments, and to run numerical sanity checks of Lemmas 9 and 12. Parts of the appendix are adapted from an earlier draft whose language was edited with OpenAI Codex. [Authors: state the extent of human review here, e.g. that all statements and proofs were verified and the text rewritten by the authors, who take full responsibility for the content.]

References

- 1 Noga Alon, Thomas F. Bloom, W. T. Gowers, Daniel Litt, Will Sawin, Arul Shankar, Jacob Tsimmerman, Victor Wang, and Melanie Matchett Wood. Remarks on the disproof of the unit distance conjecture, 2026. URL: <https://doi.org/10.48550/arXiv.2605.20695>, arXiv:2605.20695.
- 2 Noga Alon, Matija Bucić, and Lisa Sauermann. Unit and distinct distances in typical norms. *Geometric and Functional Analysis*, 35(1):1–42, 2025. URL: <https://arxiv.org/abs/2302.09058>, arXiv:2302.09058.

- 3 Hilal Asi and John C. Duchi. Instance-optimality in differential privacy via approximate inverse sensitivity mechanisms. In *Advances in Neural Information Processing Systems 33 (NeurIPS)*, pages 14106–14117, 2020. URL: <https://proceedings.neurips.cc/paper/2020/hash/a267f936e54d7c10a2bb70dbe6ad7a89-Abstract.html>.
- 4 Jeremiah Blocki, Avrim Blum, Anupam Datta, and Or Sheffet. Differentially private data analysis of social networks via restricted sensitivity. In *Proceedings of the 4th Conference on Innovations in Theoretical Computer Science (ITCS)*, pages 87–96. ACM, 2013. doi:10.1145/2422436.2422449.
- 5 Peter Brass, William O. J. Moser, and János Pach. *Research Problems in Discrete Geometry*. Springer, 2005. doi:10.1007/0-387-29929-7.
- 6 Hongjie Chen, Jingqiu Ding, Tommaso d’Orsi, Yiding Hua, Chih-Hung Liu, and David Steurer. Private graphon estimation via sum-of-squares. In *Proceedings of the 56th Annual ACM Symposium on Theory of Computing (STOC)*, pages 172–182. ACM, 2024. doi:10.1145/3618260.3649643.
- 7 Kenneth L. Clarkson, Herbert Edelsbrunner, Leonidas J. Guibas, Micha Sharir, and Emo Welzl. Combinatorial complexity bounds for arrangements of curves and spheres. *Discrete & Computational Geometry*, 5:99–160, 1990. doi:10.1007/BF02187783.
- 8 Rachel Cummings and David Durfee. Individual sensitivity preprocessing for data privacy. In *Proceedings of the 2020 ACM-SIAM Symposium on Discrete Algorithms (SODA)*, pages 528–547. SIAM, 2020. doi:10.1137/1.9781611975994.32.
- 9 Wei Dong, Juanru Fang, Ke Yi, Yuchao Tao, and Ashwin Machanavajjhala. R2T: Instance-optimal truncation for differentially private query evaluation with foreign keys. In *Proceedings of the 2022 International Conference on Management of Data (SIGMOD)*. ACM, 2022.
- 10 Cynthia Dwork, Frank McSherry, Kobbi Nissim, and Adam D. Smith. Calibrating noise to sensitivity in private data analysis. In *Theory of Cryptography*, volume 3876 of *Lecture Notes in Computer Science*, pages 265–284. Springer, 2006. doi:10.1007/11681878_14.
- 11 Cynthia Dwork and Aaron Roth. The algorithmic foundations of differential privacy. *Foundations and Trends in Theoretical Computer Science*, 9(3–4):211–407, 2014. doi:10.1561/04000000042.
- 12 Michael T. M. Emmerich. Optimizing explicit unit-distance lower-bound certificates, 2026. URL: <https://doi.org/10.48550/arXiv.2606.03419>, arXiv:2606.03419.
- 13 Paul Erdős. On sets of distances of n points. *American Mathematical Monthly*, 53(5):248–250, 1946. URL: <https://www.jstor.org/stable/2305092>.
- 14 Juanru Fang, Wei Dong, and Ke Yi. Shifted inverse: A general mechanism for monotonic functions under user differential privacy. In *Proceedings of the 2022 ACM SIGSAC Conference on Computer and Communications Security (CCS)*, pages 1009–1022. ACM, 2022. doi:10.1145/3548606.3560567.
- 15 Harold N. Gabow. Data structures for weighted matching and extensions to b -matching and f -factors. *ACM Transactions on Algorithms*, 14(3):39:1–39:80, 2018. doi:10.1145/3183369.
- 16 G. H. Hardy and E. M. Wright. *An Introduction to the Theory of Numbers*. Oxford University Press, 6 edition, 2008. Revised by D. R. Heath-Brown and J. H. Silverman.
- 17 Palak Jain, Adam D. Smith, and Connor Wagaman. Time-aware projections: Truly node-private graph statistics under continual observation. In *2024 IEEE Symposium on Security and Privacy (SP)*, pages 127–145. IEEE, 2024. URL: <https://doi.org/10.48550/arXiv.2403.04630>, arXiv:2403.04630.
- 18 Iden Kalemaj, Sofya Raskhodnikova, Adam D. Smith, and Charalampos E. Tsourakakis. Node-differentially private estimation of the number of connected components. In *Proceedings of the 42nd ACM SIGMOD-SIGACT-SIGAI Symposium on Principles of Database Systems (PODS)*, pages 183–194. ACM, 2023. doi:10.1145/3584372.3588671.
- 19 Shiva Prasad Kasiviswanathan, Kobbi Nissim, Sofya Raskhodnikova, and Adam D. Smith. Analyzing graphs with node differential privacy. In *Theory of Cryptography*, volume 7785 of *Lecture*

- Notes in Computer Science, pages 457–476. Springer, 2013. doi:10.1007/978-3-642-36594-2_26.
- 20 OpenAI. An OpenAI model has disproved a central conjecture in discrete geometry, 2026. URL: <https://openai.com/index/model-disproves-discrete-geometry-conjecture/>.
- 21 János Pach and Micha Sharir. On the number of incidences between points and curves. *Combinatorics, Probability and Computing*, 7(1):121–127, 1998. doi:10.1017/S0963548397003192.
- 22 Sofya Raskhodnikova, Adam Smith, Connor Wagaman, and Anatoly Zavyalov. Local node differential privacy, 2026. URL: <https://arxiv.org/abs/2602.15802>, arXiv:2602.15802.
- 23 Sofya Raskhodnikova and Adam D. Smith. Lipschitz extensions for node-private graph statistics and the generalized exponential mechanism. In *57th Annual IEEE Symposium on Foundations of Computer Science (FOCS)*, pages 495–504. IEEE Computer Society, 2016. doi:10.1109/FOCS.2016.60.
- 24 Will Sawin. An explicit lower bound for the unit distance problem, 2026. URL: <https://doi.org/10.48550/arXiv.2605.20579>, arXiv:2605.20579.
- 25 Alexander Schrijver. *Combinatorial Optimization: Polyhedra and Efficiency*, volume 24 of *Algorithms and Combinatorics*. Springer, 2003.
- 26 Joel Spencer, Endre Szemerédi, and William T. Trotter, Jr. Unit distances in the Euclidean plane. In Béla Bollobás, editor, *Graph Theory and Combinatorics*, pages 293–303. Academic Press, 1984.
- 27 Konrad J. Swanepoel. Combinatorial distance geometry in normed spaces. In *New Trends in Intuitive Geometry*, volume 27 of *Bolyai Society Mathematical Studies*, pages 407–458. Springer, 2018. URL: <https://arxiv.org/abs/1702.00066>, arXiv:1702.00066.
- 28 László A. Székely. Crossing numbers and hard Erdős problems in discrete geometry. *Combinatorics, Probability and Computing*, 6(3):353–358, 1997.
- 29 Endre Szemerédi and William T. Trotter, Jr. Extremal problems in discrete geometry. *Combinatorica*, 3:381–392, 1983. doi:10.1007/BF02579194.
- 30 Pavel Valtr. Strictly convex norms allowing many unit distances and related touching questions, 2005. Manuscript. The construction is reproduced in the survey by K. J. Swanepoel, *Combinatorial distance geometry in normed spaces*, 2018.
- 31 Joshua Zahl. Breaking the $3/2$ barrier for unit distances in three dimensions. *International Mathematics Research Notices*, 2019(20):6235–6284, 2019. arXiv:1706.05118.

Appendix

A Surrogate statistics

The flow statistic. Kasiviswanathan, Nissim, Raskhodnikova and Smith [19] define the following network for a graph $G = (V, E)$ and a capacity $D \geq 0$. Take a source, a sink, and two copies v_ℓ, v_r of each vertex. Add arcs of capacity D from the source to each v_ℓ and from each v_r to the sink, and for each edge uv add arcs (u_ℓ, v_r) and (v_ℓ, u_r) of capacity 1. Let $\phi_D(G)$ be the maximum flow value. The definition makes sense for real D .

► **Proposition 14.** *For every graph G and real $D \geq 0$, $\phi_D(G)/2 = g_D(G)$.*

Proof. Given a feasible x for g_D , send x_{uv} units along both arcs of each edge. This is a feasible flow of value $2 \sum_e x_e$. Conversely, from a feasible flow y set $x_{uv} = (y_{u_\ell v_r} + y_{v_\ell u_r})/2$. Then $0 \leq x_e \leq 1$, and $\sum_{e \ni v} x_e$ is half the flow through v_ℓ plus half the flow through v_r , which is at most D . The objective is half the flow value. ◀

Hence Lemma 9 is a property of the statistic of [19]; our only additions are the explicit fractional solution in part (b), which gives the bias bound for every real D , and the remark that real thresholds are needed in Theorem 1.

The integral statistic. For an integer $D \geq 0$ let $f_D(G)$ be the maximum number of edges in a subgraph of G with maximum degree at most D .

► **Proposition 15.** *For integer $D \geq 0$: (a) $|f_D(G) - f_D(G - v)| \leq D$; (b) $|E| - \text{tail}_D(G) \leq f_D(G) \leq g_D(G) \leq |E|$, and both inner inequalities can be strict; (c) f_D and g_D can be computed in polynomial time.*

Proof. (a) An optimal subgraph for G has at most D edges at v ; removing them gives a feasible subgraph of $G - v$. Every feasible subgraph of $G - v$ is feasible for G . (b) Repeatedly delete an edge at a vertex of current degree above D ; each deletion lowers $\sum_v (\deg v - D)_+$ by at least one, so at most $\text{tail}_D(G)$ edges are deleted. Integral solutions are feasible for g_D . For a triangle and $D = 1$, $f_1 = 1 < 3/2 = g_1$; for a star with more than D leaves, $g_D = D < |E|$. (c) g_D is a maximum flow; f_D is a maximum-cardinality simple b -matching with $b \equiv D$ [25, 15]. ◀

With $D' = \lceil \Lambda_k(n)/k \rceil$ in place of D in the proof of Theorem 1, the bias is at most $\text{tail}_{D'} \leq \text{tail}_D \leq \Lambda_k(n)$ and the noise is at most $(D + 1)/\varepsilon \leq \Lambda_k(n) + 1/\varepsilon$, so $f_{D'}$ gives error at most $2\Lambda_k(n) + 1/\varepsilon$.

B The adaptive algorithm

In this appendix $\log$ is the natural logarithm. Let $\mathcal{D} = \{1, 2, 4, \dots, 2^{K-1}\}$ with $K = 1 + \lceil \log_2 n \rceil$, so $\max \mathcal{D} \geq n - 1$. Fix $\beta \in (0, \frac{1}{2})$, split $\varepsilon = \varepsilon_1 + \varepsilon_2$ with $\varepsilon_1 = \varepsilon_2 = \varepsilon/2$, and let $a = \log(2/\beta)$.

The selection tool. The generalized exponential mechanism [23] of Raskhodnikova and Smith takes K candidates with score functions $q_1, \dots, q_K$, which are to be minimized; q_i has sensitivity at most Δ_i . The mechanism is ε' -differentially private, and for every $\beta' \in (0, 1)$, with probability at least $1 - \beta'$ its output $\hat{i}$ satisfies

$$q_i \leq \min_i (q_i + 4\Delta_i \log(K/\beta')/\varepsilon').$$

We use only this guarantee.

Algorithm. Run the generalized exponential mechanism with privacy parameter ε_1 on the scores $q_D(P) = -g_D(P) + aD/\varepsilon_2$, $D \in \mathcal{D}$, where score q_D is declared to have sensitivity D ; let $\hat{D}$ be the output. Release $g_{\hat{D}}(P) + Z$ with Z Laplace of scale $\hat{D}/\varepsilon_2$.

► **Lemma 16.** *The algorithm is ε -node-private. With probability at least $1 - \beta$ its error is at most $\min_{D \in \mathcal{D}} \{\text{tail}_D(P) + aD/\varepsilon_2 + 4D \log(2K/\beta)/\varepsilon_1\}$.*

Proof. Each score q_D has sensitivity at most D by Lemma 9(a), since aD/ε_2 does not depend on P . The selection step is ε_1 -node-private by the guarantee above, the release step is ε_2 -node-private given $\hat{D}$, and the composition is ε -node-private. By the guarantee above with $\beta' = \beta/2$, with probability at least $1 - \beta/2$, $q_{\hat{D}}(P) \leq \min_D \{q_D(P) + 4D \log(2K/\beta)/\varepsilon_1\}$. With probability at least $1 - \beta/2$, $|Z| \leq a\hat{D}/\varepsilon_2$. On both events,

$$|g_{\hat{D}}(P) + Z - U(P)| \leq U(P) - g_{\hat{D}}(P) + a\hat{D}/\varepsilon_2 = U(P) + q_{\hat{D}}(P),$$

and $U(P) + q_D(P) \leq \text{tail}_D(P) + aD/\varepsilon_2$ by Lemma 9(b). ◀

► **Lemma 17.** *Let $d_{(1)} \geq d_{(2)} \geq \dots$ be the degrees of $G(P)$, padded with zeros. For every integer $t \geq 1$, $\min_{D \geq 0} \{\text{tail}_D(P) + tD\} = S_t(P)$, and $\min_{D \in \mathcal{D}} \{\text{tail}_D(P) + tD\} \leq 2S_t(P) + t$.*

Proof. For every $D \geq 0$, $\text{tail}_D(P) + tD \geq \sum_{i \leq t} (d_{(i)} - D) + tD = S_t(P)$, with equality at $D = d_{(t+1)}$. For the grid, if $d_{(t+1)} = 0$ take $D = 1$: only the top t points have positive degree, so $\text{tail}_1(P) \leq S_t(P)$. Otherwise take $D \in \mathcal{D}$ with $d_{(t+1)} \leq D \leq 2d_{(t+1)}$; then $\text{tail}_D(P) + tD \leq S_t(P) - t d_{(t+1)} + 2t d_{(t+1)} \leq 2S_t(P)$. ◀

Proof of Theorem 13. By Lemma 16, the error is at most $\min_{D \in \mathcal{D}} \{\text{tail}_D(P) + t'D\}$, where $t' = \lceil 2a/\varepsilon + 8 \log(2K/\beta)/\varepsilon \rceil = O(t)$. By Lemma 17 this is at most $2S_{t'}(P) + t'$, and $S_{t'}(P) \leq \lceil t'/t \rceil S_t(P)$ because $S_j(P)/j$ is non-increasing in j . Hence the error is $O(S_t(P) + t)$. For the second statement let $k = \lceil 1/\varepsilon \rceil$. Then $S_t(P) \leq \Lambda_{\max\{t,k\}}(n) \leq (t/k + 1)\Lambda_k(n)$ by Lemma 8, and $t/k \leq \log(K/\beta) + 1$. ◀

C Deferred geometric proofs

► **Lemma 18.** *Let $\mathbb{S}$ be the unit circle of a strictly convex norm. Two distinct translates of $\mathbb{S}$ meet in at most two points, and two distinct points lie on at most two common translates of $\mathbb{S}$.*

Proof. Let K be the unit ball and $v \neq 0$. Choose coordinates (s, τ) with τ measured along v . Each line $\{s = s_0\}$ that meets the interior of K meets K in a segment $[\tau_-(s_0), \tau_+(s_0)]$, where τ_- is convex and τ_+ is concave; at the two extreme values of s the intersection is a single point, by strict convexity. The boundary of $K + v$ on the same line is $\{\tau_-(s_0) + |v|, \tau_+(s_0) + |v|\}$. Since $\tau_- \leq \tau_+$ and $|v| > 0$, a common boundary point must satisfy $\tau_+(s_0) = \tau_-(s_0) + |v|$, that is, $h(s_0) = |v|$ for $h = \tau_+ - \tau_-$. Strict convexity of K means that the graphs of $\tau_{\pm}$ contain no segment, so h is strictly concave, and it vanishes at the extreme values of s . A strictly concave function takes a positive value at most twice, which proves the first claim. For the second, $p, q \in c + \mathbb{S}$ means $c \in (p + \mathbb{S}) \cap (q + \mathbb{S})$ because $\mathbb{S} = -\mathbb{S}$, and these two translates are distinct. ◀

► **Lemma 19.** *There is an absolute constant C_0 such that for every strictly convex norm, every finite set P and every finite family Γ of distinct unit circles, $I(P, \Gamma) \leq C_0(|P|^{2/3}|\Gamma|^{2/3} + |P| + |\Gamma|)$.*

Proof. This is the argument of Székely [28]. Let $N = |P|$ and $m = |\Gamma|$. Draw a multigraph H on P : for every $\gamma \in \Gamma$ containing $j \geq 2$ points of P , join consecutive points along γ by the j arcs between them. Then H has at least $I(P, \Gamma) - m$ edges. By Lemma 18, two points are joined by arcs of at most two circles, and each circle contributes at most two such arcs, so every edge has multiplicity at most 4. Two arcs of the same circle do not cross, and two circles meet in at most two points, so the drawing has at most m^2 crossings. The crossing lemma for multigraphs with bounded multiplicity [28] states that a multigraph with N vertices, e edges and multiplicity at most μ has crossing number at least $c_1 e^3 / (\mu N^2)$ whenever $e \geq c_2 \mu N$, for absolute constants $c_1, c_2 > 0$. Hence either $I(P, \Gamma) - m < 4c_2 N$ or $m^2 \geq c_1 (I(P, \Gamma) - m)^3 / (4N^2)$, and in both cases the claim follows. $\blacktriangleleft$

Completion of the proof of Theorem 3(3). Let $n \geq 8$ and $m = \lfloor n/2 \rfloor \geq 4$, so $n/3 \leq m$ and $\sqrt{n} \leq m$.

Case $\sqrt{n} \leq k \leq m$. Put $A_0 = (m/2)^{2/3} k^{-1/3}$ and $B_0 = k^{2/3} (m/2)^{-1/3}$, so that $2A_0^2 B_0 = m$ and $A_0 B_0^2 = k$. Since $k \leq m$ and $m \geq 4$ we have $(m/2)^2 \geq m \geq k$, so $A_0 \geq 1$; since $k^2 \geq n \geq m/2$ we have $B_0 \geq 1$. Let $A = \lfloor A_0 \rfloor \geq A_0/2$ and $B = \lfloor B_0 \rfloor \geq B_0/2$. Then $2A^2 B \leq m$, $AB^2 \leq k$, and $A^2 B^2 \geq A_0^2 B_0^2 / 16 = (mk/2)^{2/3} / 16 \geq (nk/6)^{2/3} / 16 \geq n^{2/3} k^{2/3} / 53$. Let Q be the union of the $2A^2 B$ points and the AB^2 centres of Lemma 12; it has at most $m + k \leq n$ elements. The $AB^2 \leq k$ centres have total degree in $G(Q)$ at least $A^2 B^2$, so $\Lambda_k(n) \geq S_k(Q) \geq n^{2/3} k^{2/3} / 53$.

Case $m < k \leq n$. By monotonicity in k and the previous case with $k = m$, $\Lambda_k(n) \geq \Lambda_m(n) \geq n^{2/3} m^{2/3} / 53 \geq n^{2/3} k^{2/3} / 111$, since $m \geq n/3 \geq k/3$.

Case $k < \sqrt{n}$. Here $\Lambda_k(n) \geq n - 1 \geq n/2 \geq n^{2/3} k^{2/3} / 2$.

In all cases $\Lambda_k(n) \geq \max\{n^{2/3} k^{2/3} / 111, n/2\} \geq (n^{2/3} k^{2/3} + n) / 222$, so (3) holds with $c = 1/222$. $\blacktriangleleft$

D A class defined by rich points

In this appendix the norm is Euclidean. For $j \geq 1$ let $R_j(P)$ be the number of points of P of degree at least j . For $C \geq 1$ and $D_0 \geq 1$, call P (C, D_0) -regular if $R_j(P) \leq C|P|^2/j^3$ for every integer $j > D_0$. Deleting a point can violate this condition, so the class is not closed under deletion, and Theorem 1 does not apply. The bound $R_j(P) = O(n^2/j^3 + n/j)$ holds for every set by the incidence bound, so the condition removes exactly the term n/j coming from many stars of degree about j .

Proposition 20. *If P is (C, D_0) -regular, $|P| = n$, and $D \geq D_0$ is an integer, then $\text{tail}_D(P) = \sum_{j>D} R_j(P) = O(Cn^2/D^2)$. Hence the algorithm $g_D(P) + \text{Lap}(D/\varepsilon)$ with $D \asymp (C\varepsilon n^2)^{1/3} \geq D_0$ has expected error $O(C^{1/3} n^{2/3} \varepsilon^{-2/3})$, and the algorithm of Theorem 13 has error $O(C^{1/3} n^{2/3} (\log(K/\beta)/\varepsilon)^{2/3})$ with probability $1 - \beta$ whenever $D_0 \leq (Cn^2 \varepsilon / \log(K/\beta))^{1/3}$.*

Proof. $\text{tail}_D(P) = \sum_{j>D} R_j(P) \leq Cn^2 \sum_{j>D} j^{-3} = O(Cn^2/D^2)$. The first algorithm is private on all inputs by Lemma 9(a), and its error is $O(D/\varepsilon + Cn^2/D^2)$. For the second, Lemma 16 bounds the error by $O(\min_{D \in \mathcal{D}, D \geq D_0} \{Cn^2/D^2 + D \log(K/\beta)/\varepsilon\})$; optimise over D . $\blacktriangleleft$

Proposition 21. *There are absolute constants $a, b > 0$ and n_0 such that for $n \geq n_0$, $1 \leq C \leq n$, $D_0 \geq 1$ and $aC^{1/2} n^{-1/2} \leq \varepsilon \leq 1$, every ε -node-private algorithm has expected error at least $bC^{1/3} n^{2/3} \varepsilon^{-2/3}$ on some (C, D_0) -regular set of at most n points.*

Proof. Let $h = \max\{1, \lfloor 1/(4\varepsilon) \rfloor\}$, so $1/(8\varepsilon) \leq h \leq 1/\varepsilon$. Choose a large and $\alpha \leq \frac{1}{2}$ small with $2\alpha a^{-2/3} \leq \frac{1}{2}$, and set $L = \lfloor \alpha(Cn^2/h)^{1/3} \rfloor$. Since $(Cn^2/h)^{1/3} \geq (Cn^2\varepsilon)^{1/3} \geq a^{1/3}C^{1/2}n^{1/2}$, for $n \geq n_0$ we get $L \geq \frac{\alpha}{2}(Cn^2/h)^{1/3} \geq 2$ and $h(L+1) \leq 2\alpha C^{1/3}n^{2/3}\varepsilon^{-2/3} \leq n/2$. Take h far-apart gadgets, each consisting of a potential centre c_i and L points on its unit circle, with angles chosen so that no two of these L points are at unit distance, together with $n - h(L+1)$ isolated padding points. Let P_i contain all non-centre points and $c_1, \dots, c_i$. Then $|P_h| = n$, $|P_i| \geq n/2$, consecutive sets are node neighbours, and $U(P_h) - U(P_0) = hL$. For $2 \leq j \leq L$ the points of degree at least j are the inserted centres, so $R_j(P_i) \leq h \leq C|P_i|^2/L^3 \leq C|P_i|^2/j^3$ by the choice of α , and $R_j(P_i) = 0$ for $j > L$. Thus every P_i is (C, D_0) -regular. Since $h\varepsilon \leq 1$, Lemma 7 gives error at least $hL/(2(1+e)) = \Omega(C^{1/3}n^{2/3}h^{2/3}) = \Omega(C^{1/3}n^{2/3}\varepsilon^{-2/3})$. ◀

The class is useful as a conditional statement: it describes inputs on which the rich points are spread over many degree levels. We do not know a natural family of point sets with many unit distances that satisfies the condition and is not built from stars, and we leave this as a question.

E Grid points and separated sets

Grid points. Let $P \subset \frac{1}{q}\mathbb{Z}^2$ and let the target distance r satisfy $qr \in \mathbb{Z}$. Points of $\frac{1}{q}\mathbb{Z}^2$ at distance r from a grid point correspond to integer solutions of $x^2 + y^2 = (qr)^2$, so every degree is at most $\Delta = r_2((qr)^2)$, the number of such representations, and $\Delta \leq (qr)^{O(1/\log \log(qr))}$ [16]. Grid point sets form a family closed under deletion, with $\Lambda_k \leq k\Delta$. A lattice point together with the Δ grid points on its circle shows $\Lambda_k \geq k\Delta$ when $n \geq k(\Delta + 1)$ (use k far-apart copies). By Theorem 1 the error is $\Theta(\Delta/\varepsilon)$ when $n \geq \lceil 1/\varepsilon \rceil(\Delta + 1)$; the Laplace mechanism applied to U itself achieves the upper bound. If arbitrary inputs are first rounded to a fixed public grid, neighbours are mapped to equal or neighbouring grid sets, so the rounded statistic inherits privacy.

Separated sets and annuli. Fix $\delta > 0$, $r > 0$, $t \geq 0$, and let two points be adjacent when their distance lies in $[r - t, r + t]$. Consider the family of δ -separated sets, in which every two points have distance at least δ ; it is closed under deletion. The disks of radius $\delta/2$ around the neighbours of a point p are disjoint and lie in the annulus around p with radii $\max\{0, r - t - \delta/2\}$ and $r + t + \delta/2$. If the inner radius is positive, this annulus has area $2\pi r(2t + \delta)$; otherwise $r \leq t + \delta/2$ and the area is at most $\pi(2t + \delta)^2$. In both cases the area is at most $\pi(2t + \delta)(2r + 2t + \delta)$, so every degree is at most

$$\Delta_{\text{sep}} = \frac{4(2t + \delta)(2r + 2t + \delta)}{\delta^2} = O\left(\frac{rt + t^2}{\delta^2} + \frac{r + t}{\delta} + 1\right).$$

For $t \leq r$ this is $O(rt/\delta^2 + r/\delta + 1)$. For $t > r$ the term t^2/δ^2 cannot be dropped: then every point within distance $r + t$ of p is a neighbour of p , and a disk of radius $t \geq \delta$ contains $\Omega(t^2/\delta^2)$ points of a δ -separated grid. The algorithm $g_D(P) + \text{Lap}(D/\varepsilon)$ with $D = \Delta_{\text{sep}}$ is ε -node-private on *all* finite sets by Lemma 9(a); on δ -separated sets every degree is at most D , so $g_D(P) = U(P)$ by Lemma 9(b) and the only error is the noise, $\Delta_{\text{sep}}/\varepsilon$ in expectation. The same argument applies to threshold graphs (distance at most r), with degree $O(r^2/\delta^2 + 1)$.

Several distance bins. For a histogram of distances with bins $b_1, \dots, b_L$, adding a point changes all bins at once. A simple private release splits $\varepsilon = \varepsilon_1 + \dots + \varepsilon_L$ and runs the one-bin algorithm with parameter ε_i in bin i .